

Offline speech-to-text (ASR)

Offline speech-to-text (ASR)

1. Introduction
 - 1.1 What is "ASR"?
 - 1.2 Implementation Principles
 1. **Acoustic Model**
 2. **Language Model**
 3. **Pronunciation Dictionary**
 4. **Decoder**
 5. **End-to-End ASR**
2. Code Analysis
 - Key Code
 1. Speech Processing and Recognition Core ([largetmodel/largetmodel/asr.py](#))
 2. VAD Intelligent Recording ([largetmodel/largetmodel/asr.py](#))
 - Code Analysis
3. Practical Operations
 - 3.1 Configuring Offline ASR
 - 3.2 Starting and Testing the Function

1. Introduction

1.1 What is "ASR"?

ASR (Automatic Speech Recognition) is a technology that converts human speech signals into text. It is widely used in intelligent assistants, voice command control, telephone customer service automation, and real-time subtitle generation. The goal of ASR is to enable machines to "understand" human speech and convert it into a form that computers can process and understand.

1.2 Implementation Principles

The implementation of an ASR system relies primarily on the following key technical components:

1. Acoustic Model

- The acoustic model is responsible for converting the input audio signal into phonemes or subword units. This typically involves a feature extraction step, such as using Mel-Frequency Cepstral Coefficients (MFCCs) or filter banks to represent the audio signal.
- These features are then fed into a deep neural network (DNN), convolutional neural network (CNN), recurrent neural network (RNN), or the more advanced Transformer architecture for training to learn how to map audio features to corresponding phonemes or text.

2. Language Model

- A language model predicts the most likely next word given the previous context, thereby improving recognition accuracy. It is trained on large amounts of text data to understand which word sequences are most likely to occur.
- Common language models include n-gram models, RNN-based language models (LMs), and the more recently popular Transformer-based language models.

3. Pronunciation Dictionary

- A pronunciation dictionary provides a mapping between words and their corresponding pronunciations. This is crucial for connecting the acoustic model and language model, as it allows the system to understand and match the sounds it hears based on known pronunciation rules.

4. Decoder

- The decoder's task is to find the most likely word sequence as output given the acoustic model, language model, and pronunciation dictionary. This process typically involves complex search algorithms, such as the Viterbi algorithm or graph-based search methods, to find the optimal path.

5. End-to-End ASR

- With the advancement of deep learning, end-to-end ASR systems have emerged. These systems attempt to learn text output directly from raw audio signals, without requiring explicit acoustic models, separate pronunciation lexicons, and language models. These systems are often based on sequence-to-sequence (Seq2Seq) frameworks, using, for example, attention mechanisms or the Transformer architecture, significantly simplifying the complexity of traditional ASR systems.

In general, modern ASR systems achieve efficient and accurate speech-to-text conversion by combining the aforementioned components and leveraging large datasets and powerful computing resources for training. With technological advancements, the performance of ASR systems continues to improve, and their application scenarios are becoming increasingly broad.

2. Code Analysis

Key Code

1. Speech Processing and Recognition Core (`largemodel/largemodel/asr.py`)

```
# From largemodel/largemodel/asr.py
def kws_handler(self)->None:
    if self.stop_event.is_set():
        return

    if self.listen_for_speech(self.mic_index):
        asr_text = self.ASR_conversion(self.user_speechdir) # 进行 ASR 转换 /
        Perform ASR conversion
        if asr_text == 'error': # 检查 ASR 结果长度是否小于4个字符 / Check if ASR
            result length is less than 4 characters
            self.get_logger().warn("I still don't understand what you mean.
            Please try again")
            playsound(self.audio_dict[self.error_response]) # 错误响应 / Error
            response
        else:
            self.get_logger().info(asr_text)
            self.get_logger().info("okay😊, let me think for a moment...")
            self.asr_pub_result(asr_text) # 发布 ASR结果 / Publish ASR result

    else:
```

```

        return

def ASR_conversion(self, input_file:str)->str:
    if self.use_online_asr:
        result=self.modelinterface.online_asr(input_file)
        if result[0] == 'ok' and len(result[1]) > 4:
            return result[1]
        else:
            self.get_logger().error(f'ASR Error:{result[1]}') # ASR error.
            return 'error'
    else:
        result=self.modelinterface.SenseVoiceSmall_ASR(input_file)
        if result[0] == 'ok' and len(result[1]) > 4:
            return result[1]
        else:
            self.get_logger().error(f'ASR Error:{result[1]}') # ASR error.
            return 'error'

```

2. VAD Intelligent Recording (largemodel/largemodel/asr.py)

```

# From largemodel/largemodel/asr.py
def listen_for_speech(self,mic_index=0):
    p = pyaudio.PyAudio() # Create PyAudio instance. / 创建PyAudio实例。
    audio_buffer = [] # Store audio data. / 存储音频数据。
    silence_counter = 0 # Silence counter. / 静音计数器。
    MAX_SILENCE_FRAMES = 90 # 30帧*30ms=900ms静音后停止 / Stop after 900ms of
    silence (30 frames * 30ms)
    speaking = False # Flag indicating speech activity. / 语音活动标志。
    frame_counter = 0 # Frame counter. / 计数器。
    stream_kwargs = {
        'format': pyaudio.paInt16,
        'channels': 1,
        'rate': self.sample_rate,
        'input': True,
        'frames_per_buffer': self.frame_bytes,
    }
    if mic_index != 0:
        stream_kwargs['input_device_index'] = mic_index

    # Prompt the user to speak via the buzzer. / 通过蜂鸣器提示用户讲话。
    self.pub_beep.publish(UInt16(data = 1))
    time.sleep(0.5)
    self.pub_beep.publish(UInt16(data = 0))

    try:
        # Open audio stream. / 打开音频流。
        stream = p.open(**stream_kwargs)
        while True:
            if self.stop_event.is_set():
                return False

            frame = stream.read(self.frame_bytes, exception_on_overflow=False)
            # Read audio data. / 读取音频数据。
            is_speech = self.vad.is_speech(frame, self.sample_rate) # VAD
            detection. / VAD检测。

            if is_speech:

```

```

        # Detected speech activity. / 检测到语音活动。
        speaking = True
        audio_buffer.append(frame)
        silence_counter = 0
    else:
        if speaking:
            # Detect silence after speech activity. / 在语音活动后检测静音。
            silence_counter += 1
            audio_buffer.append(frame) # Continue recording buffer. / 持
续记录缓冲。

            # End recording when silence duration meets the threshold. /
静音持续时间达标时结束录音。
            if silence_counter >= MAX_SILENCE_FRAMES:
                break
        frame_counter += 1
        if frame_counter % 2 == 0:
            self.get_logger().info('1' if is_speech else '-')
            # Real-time status display.
    finally:
        stream.stop_stream()
        stream.close()
        p.terminate()

    # Save valid recording (remove trailing silence). / 保存有效录音（去除尾部静音）。
    if speaking and len(audio_buffer) > 0:
        # Trim the last silent part. / 裁剪最后静音部分。
        clean_buffer = audio_buffer[:-MAX_SILENCE_FRAMES] if len(audio_buffer) >
MAX_SILENCE_FRAMES else audio_buffer

        with wave.open(self.user_speechdir, 'wb') as wf:
            wf.setnchannels(1)
            wf.setsampwidth(p.get_sample_size(pyaudio.paInt16))
            wf.setframerate(self.sample_rate)
            wf.writeframes(b''.join(clean_buffer))
        return True

```

Code Analysis

ASR (speech-to-text) functionality is provided by the `ASRNode` node (`asr.py`). This node is responsible for recording, converting, and publishing audio.

1. Audio Recording (`listen_for_speech`):

- This function uses the `pyaudio` library to capture the audio stream from the microphone.
- It integrates the `webrtcvad` library for voice activity detection (VAD). The function loops through audio frames and uses `vad.is_speech()` to determine whether each frame contains human voice.
- When speech is detected, the data is written to a buffer. Recording stops when sustained silence (defined by `MAX_SILENCE_FRAMES`) is detected.
- Finally, the audio data in the buffer is written to a `.wav` file in the `self.user_speechdir` directory.

2. Backend Selection and Execution (`ASR_conversion`):

- The `kws_handler` function calls the `ASR_conversion` function after successful recording.
- This function determines which backend implementation to call by reading the ROS parameter `use_online_asr` (a Boolean value).

- If `false`, the `self.modelinterface.SenseVoicesSmall_ASR` method is called for local recognition.
- If `true`, the `self.modelinterface.oline_asr` method is called for online recognition.
- This function passes the audio file path as a parameter to the selected method and processes the returned result.

3. Result Publishing (`asr_pub_result`):

- After `ASR_conversion` returns valid text, `kws_handler` calls the `asr_pub_result` function.
- This function wraps the text string in a `std_msgs.msg.String` message and publishes it to the `/asr` topic via the ROS publisher.

3. Practical Operations

3.1 Configuring Offline ASR

To enable offline ASR, you need to correctly configure the `yahboom.yaml` file and ensure that the local model is correctly placed.

1. Open the configuration file:

```
vim ~/yahboom_ws/src/largemodel/config/yahboom.yaml
```

2. Modify/confirm the following key configuration:

```
asr:                                     #Voice node parameters
  ros__parameters:
    # ...
    use_oline_asr: False                 # KEY: Must be set to False to
enable offline ASR
    mic_serial_port: "/dev/ttyUSB0"      # Microphone serial port alias
    mic_index: 0                         # Microphone Device Index
    language: 'en'                       # asr language, 'zh' or 'en'
    regional_setting : "international"
```

Make sure `use_oline_asr` is set to `False` to use the local model.

In the terminal, enter `ls /dev/ttyUSB*` to check if the USB device number assigned to the voice module is `USB0`. If not, replace the `0` in the configuration file with your own device number.

Select `"zh"` for Chinese and `"en"` for English.

Also, specify the path to the offline model in `large_model_interface.yaml`.

Open the file in the terminal.

```
vim ~/yahboom_ws/src/largemodel/config/large_model_interface.yaml
```


Find the configuration related to `local_asr_model`

```
# large_model_interface.yaml
## 离线语音识别 (Offline ASR)
local_asr_model:
"/root/yahboom_ws/src/largemodel/MODELS/asr/SenseVoicesSmall" # Local ASR
model path
```

3.2 Starting and Testing the Function

Note: Due to the performance limitations of the Jetson Nano, the experience is limited. We recommend trying the online version. After starting the offline speech-to-text command, the Jetson Nano will enter a suspended state while the model is loading. Please wait patiently for at least 4 minutes. If there is no response for more than 7 minutes, the Jetson Nano has crashed. Please power cycle.

Please ensure that the voice module is connected as shown below.



1. Startup Command:

Note: The startup commands for CSI cameras and USB microphone cameras are different. Please run the command corresponding to your camera.

CSI Camera

Start udp video streaming (host)

```
./start_csi.sh
```

Enter the CSI camera docker (host machine)

```
./run_csi_docker.sh
```

Start topic conversion (docker)

```
python3 ~/temp/udp_camera_publisher.py
```

View container id

```
docker ps
```

According to the container ID shown above, multiple terminals enter the same docker

```
docker exec -it container_id /bin/bash
```

Start the offline speech-to-text command (docker)

```
ros2 launch largemodel asr_only.launch.py
```

USB Camera

Enter the USB camera docker (host machine)

```
./run_usb_docker.sh
```

Start the offline speech-to-text command (docker)

```
ros2 launch largemodel asr_only.launch.py
```

2. **Test:** Say to the microphone: Hello Xiaoya, it will respond: I'm here, and then you can start talking. Finally, it will show that the recorded sound has been converted into text and printed on the terminal.

```
~/yahboom_ws$ ros2 launch largemodel asr_only.launch.py
[INFO] [launch]: All log files can be found below /home/sunrise/.ros/log/2025-08-07-15-40-45-012409-ubuntu-7665
[INFO] [launch]: Default logging verbosity is set to INFO
[INFO] [asr-1]: process started with pid [7672]
[asr-1] [INFO] [1754552481.072012346] [asr]: The asr model :SenseVoiceSmall is loaded
[asr-1] [INFO] [1754552481.517297362] [asr]: asr_node Initialization completed
[asr-1] [INFO] [1754552491.147085778] [asr]: I'm here
```